

Python 与深度学习基础

第二次作业实验报告

母子跃 PB20000010

一、在 Tiny-ImageNet 数据集上训练 Resnet 模型

注：本次作业使用的是 Bitahub 中自带的 Tiny-ImageNet 数据集，数据量较小

(1)根据 Tiny-ImageNet 图片大小（3*64*64），计算图片经过各层处理后的中间结果的大小。请列出各层的名称及输出的大小。

输入图片大小：3×64×64

第一层卷积处理：Conv1，卷积核大小为 7×7，步长为 2，填充为 3，输出特征图尺寸为：64×32×32

第二层池化处理：Maxpool，池化核大小为 3×3，步长为 2，填充为 1，输出特征图尺寸为：64×16×16

第一个残差块：layer1，包含两个基本块（BasicBlock），需要分别计算两个基本块的输出：

1) 第一个基本块：输出特征图尺寸为 64×16×16。

2) 第二个基本块：输出特征图尺寸为 64×16×16。

第二个残差块：layer2，包含两个基本块，同样需要分别计算两个基本块的输出：

1) 第一个基本块：输出特征图尺寸为 128×8×8。

2) 第二个基本块：输出特征图尺寸为 128×8×8。

第三个残差块：layer3，包含两个基本块，同样需要分别计算两个基本块的输出：

1) 第一个基本块：输出特征图尺寸为 256×4×4。

2) 第二个基本块：输出特征图尺寸为 256×4×4。

第四个残差块：layer4，包含两个基本块，同样需要分别计算两个基本块的输出：

1) 第一个基本块：输出特征图尺寸为 512×2×2。

2) 第二个基本块：输出特征图尺寸为 512×2×2。

全局平均池化层：Avgpool，池化核大小为 2×2，步长为 1，输出特征图尺寸为：512×1×1。

全连接层：Fc，将输入展开为一维向量，输出维度为 200。

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 32, 32]	9,408
BatchNorm2d-2	[-1, 64, 32, 32]	128
ReLU-3	[-1, 64, 32, 32]	0
MaxPool2d-4	[-1, 64, 16, 16]	0
Conv2d-5	[-1, 64, 16, 16]	36,864
BatchNorm2d-6	[-1, 64, 16, 16]	128
ReLU-7	[-1, 64, 16, 16]	0

Conv2d-8	[-1, 64, 16, 16]	36,864
BatchNorm2d-9	[-1, 64, 16, 16]	128
ReLU-10	[-1, 64, 16, 16]	0
BasicBlock-11	[-1, 64, 16, 16]	0
Conv2d-12	[-1, 64, 16, 16]	36,864
BatchNorm2d-13	[-1, 64, 16, 16]	128
ReLU-14	[-1, 64, 16, 16]	0
Conv2d-15	[-1, 64, 16, 16]	36,864
BatchNorm2d-16	[-1, 64, 16, 16]	128
ReLU-17	[-1, 64, 16, 16]	0
BasicBlock-18	[-1, 64, 16, 16]	0
Conv2d-19	[-1, 128, 8, 8]	73,728
BatchNorm2d-20	[-1, 128, 8, 8]	256
ReLU-21	[-1, 128, 8, 8]	0
Conv2d-22	[-1, 128, 8, 8]	147,456
BatchNorm2d-23	[-1, 128, 8, 8]	256
Conv2d-24	[-1, 128, 8, 8]	8,192
BatchNorm2d-25	[-1, 128, 8, 8]	256
ReLU-26	[-1, 128, 8, 8]	0
BasicBlock-27	[-1, 128, 8, 8]	0
Conv2d-28	[-1, 128, 8, 8]	147,456
BatchNorm2d-29	[-1, 128, 8, 8]	256
ReLU-30	[-1, 128, 8, 8]	0
Conv2d-31	[-1, 128, 8, 8]	147,456
BatchNorm2d-32	[-1, 128, 8, 8]	256
ReLU-33	[-1, 128, 8, 8]	0
BasicBlock-34	[-1, 128, 8, 8]	0
Conv2d-35	[-1, 256, 4, 4]	294,912
BatchNorm2d-36	[-1, 256, 4, 4]	512
ReLU-37	[-1, 256, 4, 4]	0
Conv2d-38	[-1, 256, 4, 4]	589,824
BatchNorm2d-39	[-1, 256, 4, 4]	512
Conv2d-40	[-1, 256, 4, 4]	32,768
BatchNorm2d-41	[-1, 256, 4, 4]	512
ReLU-42	[-1, 256, 4, 4]	0
BasicBlock-43	[-1, 256, 4, 4]	0
Conv2d-44	[-1, 256, 4, 4]	589,824
BatchNorm2d-45	[-1, 256, 4, 4]	512
ReLU-46	[-1, 256, 4, 4]	0
Conv2d-47	[-1, 256, 4, 4]	589,824
BatchNorm2d-48	[-1, 256, 4, 4]	512
ReLU-49	[-1, 256, 4, 4]	0
BasicBlock-50	[-1, 256, 4, 4]	0
Conv2d-51	[-1, 512, 2, 2]	1,179,648

BatchNorm2d-52	[-1, 512, 2, 2]	1,024
ReLU-53	[-1, 512, 2, 2]	0
Conv2d-54	[-1, 512, 2, 2]	2,359,296
BatchNorm2d-55	[-1, 512, 2, 2]	1,024
Conv2d-56	[-1, 512, 2, 2]	131,072
BatchNorm2d-57	[-1, 512, 2, 2]	1,024
ReLU-58	[-1, 512, 2, 2]	0
BasicBlock-59	[-1, 512, 2, 2]	0
Conv2d-60	[-1, 512, 2, 2]	2,359,296
BatchNorm2d-61	[-1, 512, 2, 2]	1,024
ReLU-62	[-1, 512, 2, 2]	0
Conv2d-63	[-1, 512, 2, 2]	2,359,296
BatchNorm2d-64	[-1, 512, 2, 2]	1,024
ReLU-65	[-1, 512, 2, 2]	0
BasicBlock-66	[-1, 512, 2, 2]	0
AdaptiveAvgPool2d-67	[-1, 512, 1, 1]	0
Linear-68	[-1, 200]	102,600

```

=====
Total params: 11,279,112
Trainable params: 11,279,112
Non-trainable params: 0
-----

```

```

Input size (MB): 0.05
Forward/backward pass size (MB): 5.13
Params size (MB): 43.03
Estimated Total Size (MB): 48.20

```

(2) 原代码中对应的是 ImageNet，它有 1000 类，也就是 output 有 1000 维，修改成 200 维：

```

# -----change: use function 'torch.nn.Linear()' to change the output dimension of the fully connected layer
out_features = 200 # output features for tiny-imagenet
in_features = model.fc.in_features # the original features of resnet
model.fc = torch.nn.Linear(in_features, out_features)

```

(3) 处理数据集，修改正确标签

首先，设置 traindir 变量为 Tiny-ImageNet 训练集文件夹的路径；
然后，通过打开 val_annotations.txt 文件并读取其中的标签信息，将每个图像与其对应的标签建立映射关系，并存储在字典 d 中；
接下来，根据字典 d 的内容，为每个类别创建一个新的文件夹，用于存储 Tiny-ImageNet 验证集中该类别的所有图像；
对于每张图像，从原始 Tiny-ImageNet 验证集中读取其图像数据，并将其复制到相应的新文件夹中；
最后，设置 valdir 变量为 Tiny-ImageNet 新建的验证集文件夹的路径，并进行数据归一化处理。

```

# Data Loading code
# ----- Change: We need to change the construction in valdir when using tiny-imagenet -----#
traindir = os.path.join(args.data, 'train')
# valdir = os.path.join(args.data, 'val')
f = open("/data/bitahub/Tiny-ImageNet/val/val_annotations.txt", "r") # the relative path of the file is corresponding to Local path
val_labels = f.readlines() # get Labels in val_annotations.txt to correct the original Labels
f.close()

d = {} # create a new dictionary to store the map between images and Labels
for item in val_labels:
    # split according to the format of val_annotations.txt
    image_name = item.split("\t")[0]
    image_label = item.split("\t")[1]
    if image_label not in d:
        d[image_label] = [image_name]
    else:
        d[image_label] += [image_name]
# for item in d:
#     if not os.path.exists("/data/bitahub/Tiny-ImageNet/my_val/{}/images".format(item)):
#         os.makedirs("/data/bitahub/Tiny-ImageNet/my_val/{}/images".format(item))
#         # os.makedirs() could only been used when the path is not existed
#     for num, img in enumerate(d[item]):
#         source = "/data/bitahub/Tiny-ImageNet/val/images/{}".format(img)
#         destination = "/data/bitahub/Tiny-ImageNet/my_val/{}/images/{}_{}.JPEG".format(item, item, num)
#         # using shutil.copyfile() to create new val_dir just like the train_dir
#         shutil.copyfile(source, destination)

# the dataset on bitahub is "read-only" mode, so cann't write the new folder "my_val"
for item in d:
    if not os.path.exists("/mydata/my_val/{}/images".format(item)):
        os.makedirs("/mydata/my_val/{}/images".format(item))
        # os.makedirs() could only been used when the path is not existed
    for num, img in enumerate(d[item]):
        source = "/data/bitahub/Tiny-ImageNet/val/images/{}".format(img)
        destination = "/mydata/my_val/{}/images/{}_{}.JPEG".format(item, item, num)
        # using shutil.copyfile() to create new val_dir just like the train_dir
        shutil.copyfile(source, destination)
# create new val_dir
# valdir = os.path.join(args.data, 'my_val')
valdir = '/mydata/my_val'
normalize = transforms.Normalize(mean=[0.485, 0.456, 0.406],
                                std=[0.229, 0.224, 0.225])

```

(4) 在代码中增加 `torch.utils.tensorboard` 的代码，以能在 TensorBoard 中观察训练集 Loss、训练集精度、验证集 Loss、验证集精度的变化。

首先，代码使用 `datetime.now().strftime()` 函数获取当前时间，并且将其作为 tensorboard 存储的目录。之后，将训练数据加载器传递给 `train()` 函数，获得返回的 `train_loss` 和 `train_acc` 值。同时，将验证数据加载器传递给 `validate()` 函数，获得返回的 `val_loss`、`acc1` 和 `val_acc` 值。这些值将用于绘制 scalar（标量）图，以便能够更好地理解模型的训练和验证过程。在每个 epoch 结束时，将会检查当前是否是最佳的模型，并在 `args.save_path` 路径下保存模型参数。在第 5 和 10 个 epoch 结束时，也会将模型参数保存在 `args.save_path` 路径下，以便查看模型在不同阶段的表现。最后，在 TensorBoard 中绘制 scalar 图，记录训练 loss、训练 acc、验证 loss 和验证 acc 的变化情况。在训练过程结束后，关闭 TensorBoard 的写入器 writer，输出 "Training Finished"。

```

#----- Change: no need to reshape and flip in tiny-imagenet -----#
val_loader = torch.utils.data.DataLoader(
    datasets.ImageFolder(valdir, transforms.Compose([
        # transforms.Resize(256),
        # transforms.CenterCrop(224),
        transforms.ToTensor(),
        normalize,
    ])),
    batch_size=args.batch_size, shuffle=False,
    num_workers=args.workers, pin_memory=True)

# Create Tensorboard and assign its storage dir as "datetime + name_of_web_construction"
current_time = datetime.now().strftime('%b%d_%H-%M-%S')
logdir = os.path.join('/output', 'logs', current_time + '_' + args.arch)
writer = SummaryWriter(logdir)

# Use Tensorboard to draw Graph of the net
# dummy_input = torch.rand(4, 3, 64, 64) # The corresponding input of tiny-imagenet
# writer.add_graph(model, (dummy_input,)) # model in this .py file cannot be drawn by this way,
summary(model, (3, 64, 64)) # Instead, using Lib "summary" to represent the construction

if args.evaluate:
    validate(val_loader, model, criterion, args)
    return

for epoch in range(args.start_epoch, args.epochs):
    if args.distributed:
        train_sampler.set_epoch(epoch)

    # train for one epoch
    # ----- Change: We need more returned value to draw tensorboard -----#
    train_loss, train_acc = train(train_loader, model, criterion, optimizer, epoch, args)

    # evaluate on validation set
    acc1, val_loss, val_acc = validate(val_loader, model, criterion, args)

    scheduler.step()

    # remember best acc@1 and save checkpoint
    is_best = acc1 > best_acc1
    best_acc1 = max(acc1, best_acc1)

    if not args.multiprocessing_distributed or (args.multiprocessing_distributed
        and args.rank % ngpus_per_node == 0):
        save_checkpoint({
            'epoch': epoch + 1,
            'arch': args.arch,
            'state_dict': model.state_dict(),
            'best_acc1': best_acc1,
            'optimizer' : optimizer.state_dict(),
            'scheduler' : scheduler.state_dict()
        }, is_best)

    # pick another two checkpoints for evaluation
    if epoch == 5 or epoch == 10:
        save_checkpoint({
            'epoch': epoch + 1,
            'arch': args.arch,
            'state_dict': model.state_dict(),
            'best_acc1': best_acc1,
            'optimizer' : optimizer.state_dict(),
            'scheduler' : scheduler.state_dict()
        }, is_best, "checkpoint_epoch{}.pth.tar".format(epoch))

    # when one epoch finished, save scalar in tensorboard for visibility
    writer.add_scalar('scalar/train_loss', train_loss, epoch)
    writer.add_scalar('scalar/train_acc', train_acc, epoch)
    writer.add_scalar('scalar/val_loss', val_loss, epoch)
    writer.add_scalar('scalar/val_acc', val_acc, epoch)

print("Training Finished")
writer.close()

```

在 `train()` 函数中，首先使用 `AverageMeter` 类初始化了 `batch_time`（每批数据的时间）、`data_time`（数据加载时间）、`losses`（损失）、`top1`（Top1 准确率）和 `top5`（Top5 准确率）五个度量器。之后，使用 `ProgressMeter` 类初始化了 `progress`，该类用于打印每个 `epoch` 的进度条。

在每个 epoch 开始时，将模型转换为训练模式 `model.train()`。对于每个 batch 数据，先记录数据加载所需时间 `data_time`，同时将数据 `images` 和标签 `target` 分别移动到 GPU 上（如果有）。

然后，计算模型的输出 `output` 和损失 `loss`，并通过 `accuracy()` 函数计算准确率 `acc1` 和 `acc5`。将损失值和准确率值记录在相应度量器中。

接着，清空优化器 `optimizer` 中的梯度信息 `optimizer.zero_grad()`，反向传播 `loss.backward()` 并更新模型参数 `optimizer.step()`。

最后，在每 `args.print_freq`（默认为 10）个 batch 之后，打印当前进度 `progress.display(i)`。返回 `loss` 和 `top5.avg`（Top5 准确率的平均值）两个值。

```
def train(train_loader, model, criterion, optimizer, epoch, args):
    batch_time = AverageMeter('Time', ':6.3f')
    data_time = AverageMeter('Data', ':6.3f')
    losses = AverageMeter('Loss', ':.4e')
    top1 = AverageMeter('Acc@1', ':6.2f')
    top5 = AverageMeter('Acc@5', ':6.2f')
    progress = ProgressMeter(
        len(train_loader),
        [batch_time, data_time, losses, top1, top5],
        prefix="Epoch: [{}]" .format(epoch))

    # switch to train mode
    model.train()

    end = time.time()
    for i, (images, target) in enumerate(train_loader):
        # measure data loading time
        data_time.update(time.time() - end)

        if args.gpu is not None:
            images = images.cuda(args.gpu, non_blocking=True)
        if torch.cuda.is_available():
            target = target.cuda(args.gpu, non_blocking=True)

        # compute output
        output = model(images)
        loss = criterion(output, target)

        # measure accuracy and record loss
        acc1, acc5 = accuracy(output, target, topk=(1, 5))
        losses.update(loss.item(), images.size(0))
        top1.update(acc1[0], images.size(0))
        top5.update(acc5[0], images.size(0))

        # compute gradient and do SGD step
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # measure elapsed time
        batch_time.update(time.time() - end)
        end = time.time()

        if i % args.print_freq == 0:
            progress.display(i)

    #-----Change: Add more return values-----
    return loss, top5.avg # Our aim is to boost the average c
```


在 `validate()` 函数中，依然是使用 `AverageMeter` 类分别初始化 `batch_time`（每批数据的时间）、`losses`（损失）、`top1`（Top1 准确率）和 `top5`（Top5 准确率）四个度量器。需要注意的是，在 `top1` 和 `top5` 的初始化中，我们使用了 `Summary` 枚举类来指定每次更新后是否影响其总数值，以及如何汇总每个 `batch` 中的值。

接着，在进入“测试模式”`model.eval()` 后，使用 `with torch.no_grad()` 包含 `for` 循环，避免计算图的构建和梯度更新，从而减少 GPU 存储需求和计算量。

在每个 `batch` 数据上，与 `train()` 函数类似，将数据 `images` 和标签 `target` 分别移动到 GPU 上，并计算模型的输出 `output` 和损失 `loss`。使用 `accuracy()` 函数计算准确率 `acc1` 和 `acc5`，同时在相应度量器中记录损失值和准确率值。

在计算完每个 `batch` 后，更新其他度量器 `batch_time` 和 `progress`。`progress.display(i)` 用于打印当前进度，如果 `args.evaluate` 为 `True`，则会额外打印出每张图片的预测结果和真实标签，最后调用 `progress.display_summary()` 打印所有 `epoch` 的总结信息。

最后，返回 `top1.avg`（Top1 准确率的平均值）、`loss`（损失值）和 `top5.avg`（Top5 准确率的平均值）三个值。

```
def validate(val_loader, model, criterion, args):
    batch_time = AverageMeter('Time', ':6.3f', Summary.NONE)
    losses = AverageMeter('Loss', ':.4e', Summary.NONE)
    top1 = AverageMeter('Acc@1', ':6.2f', Summary.AVERAGE)
    top5 = AverageMeter('Acc@5', ':6.2f', Summary.AVERAGE)
    progress = ProgressMeter(
        len(val_loader),
        [batch_time, losses, top1, top5],
        prefix='Test: ')

    # switch to evaluate mode
    model.eval()

    with torch.no_grad():
        end = time.time()
        for i, (images, target) in enumerate(val_loader):
            if args.gpu is not None:
                images = images.cuda(args.gpu, non_blocking=True)
            if torch.cuda.is_available():
                target = target.cuda(args.gpu, non_blocking=True)

            # compute output
            output = model(images)
            loss = criterion(output, target)

            # measure accuracy and record loss
            acc1, acc5 = accuracy(output, target, topk=(1, 5))
            losses.update(loss.item(), images.size(0))
            top1.update(acc1[0], images.size(0))
            top5.update(acc5[0], images.size(0))

            # measure elapsed time
            batch_time.update(time.time() - end)
            end = time.time()

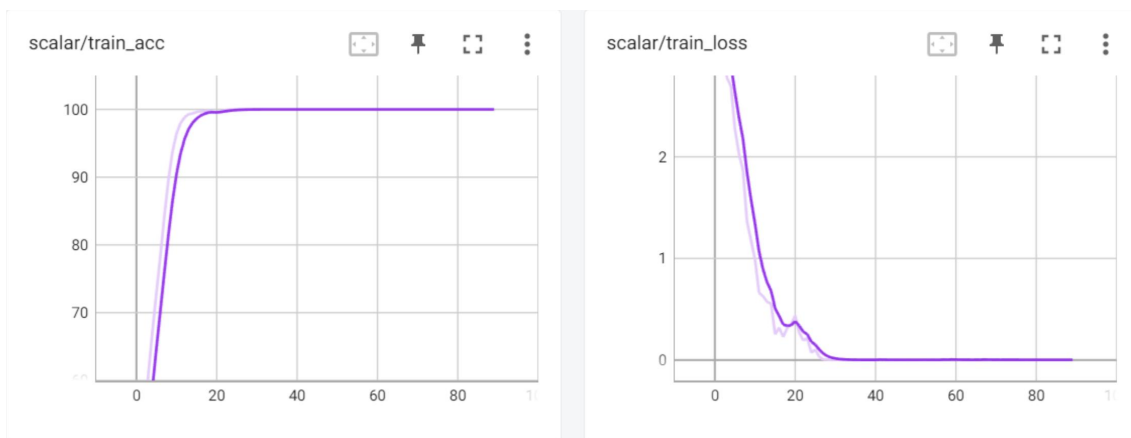
            if i % args.print_freq == 0:
                progress.display(i)

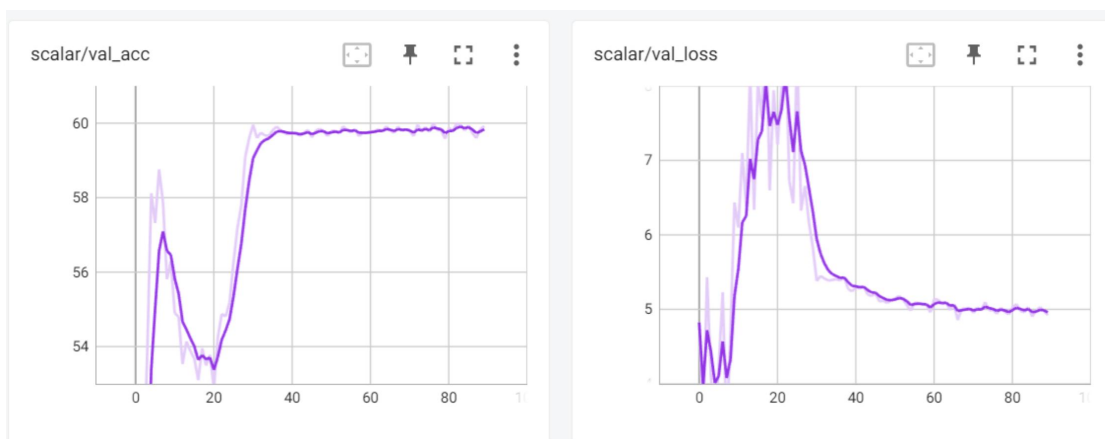
            if args.evaluate:
                # when evaluating, show more detail about each picture
                print("\rpic:{:2}  aim:{:2}  result:{:4} ({:8})".format(i, int(target), int(output.argmax()), 'correct' if int(target) == int(output.argmax()) else 'false'))
                if i == 100:
                    # no need to calculate all 10,000 pics
                    break

        progress.display_summary()

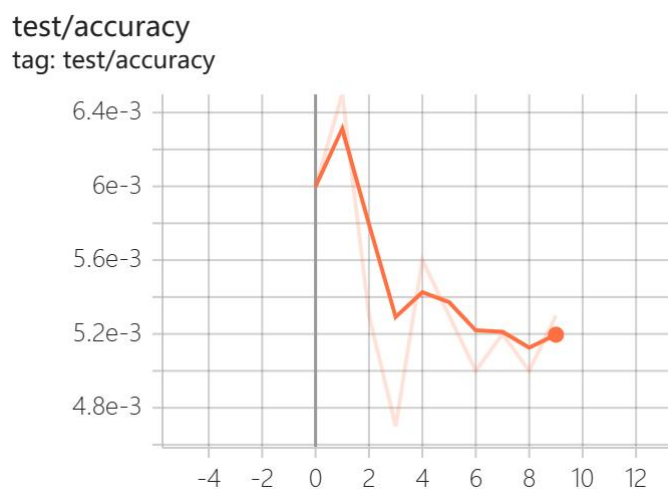
    #-----Change: Add more return values-----#
    return top1.avg, loss, top5.avg
```

(5) 对第 3 步中的曲线进行截图





需要注意的是，在前几次的训练过程中，出现了如下图中的“过拟合”现象：



(6) 分别在无 GPU、1 个 GPU、多个 GPU 环境下，重复上述过程，观察和量化评价训练速度上的差异

修改输入的 GPU 数量：

```
def main():
    args = argparse.Namespace(
        arch='resnet18',
        data_url=None,
        data_dir=None,
        workers=2,
        epochs=90,
        start_epoch=0,
        batch_size=256,
        lr=0.1,
        momentum=0.9,
        weight_decay=1e-4,
        print_freq=10,
        resume=None,
        world_size=1,
        rank=0,
        dist_url='tcp://127.0.0.1:23456',
        dist_backend='nccl',
        seed=42,
        gpu=8, # 修改gpu个数
        pretrained=True,
        data='/data/bitahub/Tiny-ImageNet',
        evaluate=True,
        multiprocessing_distributed=False
    )
```


由于训练时间过长，导致无法统计出具体的训练时间

(7) 保存 2 个训练过程中模型的 checkpoint

- 📁 checkpoint_epoch10.pth.tar
- 📁 checkpoint_epoch5.pth.tar
- 📁 checkpoint.pth.tar

(8) 运行时部分过程截图

Epoch: [14][390/391]	Time	0.037 (0.180)	Data	0.000 (0.146)	Loss	5.5408e-01 (2.9483e-01)	Acc@1	85.00 (91.01)	Acc@5	96.25 (99.39)
Test: [0/40]	Time	0.447 (0.447)	Loss	3.4266e+00 (3.4266e+00)	Acc@1	41.02 (41.02)	Acc@5	69.53 (69.53)		
Test: [10/40]	Time	0.015 (0.077)	Loss	5.4116e+00 (4.4403e+00)	Acc@1	23.44 (30.79)	Acc@5	48.83 (57.88)		
Test: [20/40]	Time	0.100 (0.063)	Loss	4.3161e+00 (4.7367e+00)	Acc@1	30.08 (28.05)	Acc@5	55.47 (54.00)		
Test: [30/40]	Time	0.015 (0.057)	Loss	5.3962e+00 (4.8278e+00)	Acc@1	19.92 (27.51)	Acc@5	47.27 (52.84)		
* Acc@1 28.070 Acc@5 53.880										
Epoch: [15][0/391]	Time	1.084 (1.084)	Data	1.054 (1.054)	Loss	1.9262e-01 (1.9262e-01)	Acc@1	94.14 (94.14)	Acc@5	100.00 (100.00)
Epoch: [15][10/391]	Time	0.043 (0.226)	Data	0.000 (0.187)	Loss	2.3546e-01 (2.4337e-01)	Acc@1	93.75 (92.37)	Acc@5	99.61 (99.72)
Epoch: [15][20/391]	Time	0.564 (0.219)	Data	0.537 (0.181)	Loss	2.6244e-01 (2.4035e-01)	Acc@1	91.02 (92.76)	Acc@5	99.61 (99.72)
Epoch: [15][30/391]	Time	0.039 (0.198)	Data	0.000 (0.160)	Loss	3.2740e-01 (2.4668e-01)	Acc@1	90.23 (92.39)	Acc@5	99.22 (99.65)
Epoch: [15][40/391]	Time	0.553 (0.197)	Data	0.523 (0.159)	Loss	2.3604e-01 (2.4709e-01)	Acc@1	92.97 (92.44)	Acc@5	100.00 (99.66)
Epoch: [15][50/391]	Time	0.045 (0.189)	Data	0.011 (0.152)	Loss	2.0222e-01 (2.4366e-01)	Acc@1	92.19 (92.56)	Acc@5	100.00 (99.63)
Epoch: [15][60/391]	Time	0.495 (0.190)	Data	0.469 (0.153)	Loss	1.6566e-01 (2.3714e-01)	Acc@1	95.31 (92.78)	Acc@5	100.00 (99.67)
Epoch: [15][70/391]	Time	0.102 (0.185)	Data	0.074 (0.149)	Loss	1.4021e-01 (2.3240e-01)	Acc@1	96.09 (92.90)	Acc@5	100.00 (99.68)
Epoch: [15][80/391]	Time	0.536 (0.186)	Data	0.510 (0.150)	Loss	1.8382e-01 (2.2665e-01)	Acc@1	93.36 (93.14)	Acc@5	99.61 (99.70)
Epoch: [15][90/391]	Time	0.043 (0.182)	Data	0.011 (0.147)	Loss	2.3303e-01 (2.2299e-01)	Acc@1	92.58 (93.32)	Acc@5	100.00 (99.73)
Epoch: [15][100/391]	Time	0.475 (0.183)	Data	0.445 (0.149)	Loss	1.5376e-01 (2.1845e-01)	Acc@1	94.92 (93.49)	Acc@5	100.00 (99.74)
Epoch: [15][110/391]	Time	0.100 (0.182)	Data	0.072 (0.147)	Loss	1.9349e-01 (2.1447e-01)	Acc@1	94.14 (93.64)	Acc@5	100.00 (99.76)
Epoch: [15][120/391]	Time	0.495 (0.182)	Data	0.465 (0.148)	Loss	2.1741e-01 (2.1163e-01)	Acc@1	92.97 (93.71)	Acc@5	100.00 (99.78)
Epoch: [15][130/391]	Time	0.122 (0.180)	Data	0.094 (0.145)	Loss	2.1334e-01 (2.0835e-01)	Acc@1	92.58 (93.84)	Acc@5	100.00 (99.78)
Epoch: [15][140/391]	Time	0.420 (0.180)	Data	0.390 (0.146)	Loss	1.8477e-01 (2.0788e-01)	Acc@1	95.31 (93.83)	Acc@5	99.61 (99.77)
Epoch: [15][150/391]	Time	0.202 (0.179)	Data	0.170 (0.145)	Loss	2.2551e-01 (2.0700e-01)	Acc@1	94.92 (93.87)	Acc@5	100.00 (99.77)
Epoch: [15][160/391]	Time	0.474 (0.180)	Data	0.447 (0.145)	Loss	1.7385e-01 (2.0579e-01)	Acc@1	94.92 (93.89)	Acc@5	100.00 (99.77)
Epoch: [15][170/391]	Time	0.260 (0.180)	Data	0.236 (0.146)	Loss	2.3954e-01 (2.0525e-01)	Acc@1	91.41 (93.90)	Acc@5	100.00 (99.77)
Epoch: [15][180/391]	Time	0.408 (0.181)	Data	0.379 (0.147)	Loss	1.4985e-01 (2.0513e-01)	Acc@1	95.31 (93.91)	Acc@5	100.00 (99.77)
Epoch: [15][190/391]	Time	0.245 (0.180)	Data	0.217 (0.146)	Loss	1.8823e-01 (2.0439e-01)	Acc@1	95.31 (93.94)	Acc@5	99.22 (99.77)
Epoch: [15][200/391]	Time	0.419 (0.181)	Data	0.392 (0.147)	Loss	2.2993e-01 (2.0544e-01)	Acc@1	94.14 (93.89)	Acc@5	99.61 (99.77)
Epoch: [15][210/391]	Time	0.191 (0.181)	Data	0.164 (0.147)	Loss	2.8236e-01 (2.0604e-01)	Acc@1	90.62 (93.86)	Acc@5	99.61 (99.77)
Epoch: [15][220/391]	Time	0.420 (0.181)	Data	0.389 (0.147)	Loss	1.9513e-01 (2.0709e-01)	Acc@1	94.92 (93.85)	Acc@5	99.61 (99.76)
Epoch: [15][230/391]	Time	0.254 (0.181)	Data	0.223 (0.147)	Loss	2.8844e-01 (2.0834e-01)	Acc@1	91.02 (93.80)	Acc@5	99.61 (99.76)
Epoch: [15][240/391]	Time	0.407 (0.181)	Data	0.374 (0.147)	Loss	2.5334e-01 (2.1014e-01)	Acc@1	91.02 (93.74)	Acc@5	99.22 (99.76)
Epoch: [22][220/391]	Time	1.216 (0.615)	Data	1.085 (0.497)	Loss	2.7916e-01 (1.8800e-01)	Acc@1	90.23 (9		
4.40)	Acc@5	99.61 (99.72)								
Epoch: [22][230/391]	Time	0.104 (0.600)	Data	0.001 (0.482)	Loss	2.2466e-01 (1.9033e-01)	Acc@1	92.19 (9		
4.33)	Acc@5	99.61 (99.70)								
Epoch: [22][240/391]	Time	0.446 (0.585)	Data	0.309 (0.467)	Loss	3.1815e-01 (1.9245e-01)	Acc@1	90.23 (9		
4.24)	Acc@5	98.83 (99.69)								
Epoch: [22][250/391]	Time	0.120 (0.569)	Data	0.001 (0.451)	Loss	2.5826e-01 (1.9396e-01)	Acc@1	91.80 (9		
4.22)	Acc@5	99.22 (99.69)								
Epoch: [22][260/391]	Time	0.473 (0.555)	Data	0.339 (0.436)	Loss	3.1189e-01 (1.9627e-01)	Acc@1	89.45 (9		
4.13)	Acc@5	99.22 (99.69)								
Epoch: [22][270/391]	Time	0.111 (0.541)	Data	0.001 (0.423)	Loss	2.2380e-01 (1.9772e-01)	Acc@1	92.58 (9		
4.08)	Acc@5	99.61 (99.69)								
Epoch: [22][280/391]	Time	1.026 (0.536)	Data	0.894 (0.418)	Loss	2.8041e-01 (2.0047e-01)	Acc@1	91.41 (9		
3.97)	Acc@5	98.83 (99.67)								
Epoch: [22][290/391]	Time	0.109 (0.531)	Data	0.001 (0.413)	Loss	2.5815e-01 (2.0272e-01)	Acc@1	91.41 (9		
3.90)	Acc@5	99.61 (99.67)								
Epoch: [22][300/391]	Time	1.212 (0.530)	Data	1.087 (0.411)	Loss	2.8415e-01 (2.0528e-01)	Acc@1	91.41 (9		
3.81)	Acc@5	98.83 (99.66)								
Epoch: [22][310/391]	Time	0.121 (0.523)	Data	0.001 (0.404)	Loss	2.6021e-01 (2.0699e-01)	Acc@1	93.36 (9		
3.77)	Acc@5	99.61 (99.65)								
Epoch: [22][320/391]	Time	0.900 (0.515)	Data	0.651 (0.396)	Loss	2.6036e-01 (2.0927e-01)	Acc@1	91.02 (9		
3.68)	Acc@5	99.22 (99.64)								
Epoch: [22][330/391]	Time	0.113 (0.505)	Data	0.001 (0.386)	Loss	2.0334e-01 (2.1194e-01)	Acc@1	95.70 (9		
3.61)	Acc@5	99.61 (99.63)								
Epoch: [22][340/391]	Time	1.226 (0.499)	Data	1.108 (0.380)	Loss	2.7873e-01 (2.1462e-01)	Acc@1	91.41 (9		
3.53)	Acc@5	99.61 (99.61)								
Epoch: [22][350/391]	Time	0.111 (0.496)	Data	0.001 (0.377)	Loss	2.6511e-01 (2.1615e-01)	Acc@1	93.36 (9		
3.49)	Acc@5	100.00 (99.61)								
Epoch: [22][360/391]	Time	1.602 (0.500)	Data	1.485 (0.381)	Loss	2.8292e-01 (2.1853e-01)	Acc@1	89.84 (9		
3.42)	Acc@5	99.61 (99.60)								
Epoch: [22][370/391]	Time	0.110 (0.497)	Data	0.001 (0.378)	Loss	4.1160e-01 (2.2198e-01)	Acc@1	88.28 (9		
3.30)	Acc@5	98.83 (99.59)								
Epoch: [22][380/391]	Time	1.701 (0.501)	Data	1.580 (0.382)	Loss	4.0028e-01 (2.2427e-01)	Acc@1	86.33 (9		
3.23)	Acc@5	99.22 (99.58)								
Epoch: [22][390/391]	Time	0.104 (0.500)	Data	0.000 (0.382)	Loss	4.7682e-01 (2.2714e-01)	Acc@1	83.75 (9		
3.14)	Acc@5	98.12 (99.57)								
Test: [0/40]	Time	1.226 (1.226)	Loss	4.5341e+00 (4.5341e+00)	Acc@1	30.08 (30.08)	Acc@5	56.64 (56.64)		
Test: [10/40]	Time	0.060 (0.191)	Loss	4.5721e+00 (4.5636e+00)	Acc@1	33.98 (30.15)	Acc@5	58.59 (55.58)		
Test: [20/40]	Time	0.073 (0.135)	Loss	4.9465e+00 (4.8074e+00)	Acc@1	28.12 (28.12)	Acc@5	49.61 (52.57)		
Test: [30/40]	Time	0.068 (0.116)	Loss	5.2082e+00 (4.8877e+00)	Acc@1	28.12 (28.23)	Acc@5	50.78 (51.80)		

Epoch: [89][150/391]	Time	0.042 (0.175)	Data	0.000 (0.140)	Loss	2.5321e-03 (2.8378e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][160/391]	Time	0.215 (0.174)	Data	0.187 (0.140)	Loss	2.4638e-03 (2.8237e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][170/391]	Time	0.039 (0.174)	Data	0.000 (0.140)	Loss	2.6011e-03 (2.8371e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][180/391]	Time	0.306 (0.173)	Data	0.278 (0.139)	Loss	2.2875e-03 (2.8340e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][190/391]	Time	0.042 (0.173)	Data	0.000 (0.138)	Loss	5.8494e-03 (2.8573e-03)	Acc@1	99.61 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][200/391]	Time	0.265 (0.172)	Data	0.236 (0.137)	Loss	2.6882e-03 (2.8647e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][210/391]	Time	0.041 (0.172)	Data	0.000 (0.137)	Loss	2.8928e-03 (2.8663e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][220/391]	Time	0.357 (0.172)	Data	0.330 (0.137)	Loss	2.6964e-03 (2.8684e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][230/391]	Time	0.042 (0.171)	Data	0.000 (0.136)	Loss	2.5581e-03 (2.8698e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][240/391]	Time	0.322 (0.171)	Data	0.295 (0.137)	Loss	2.6622e-03 (2.8719e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][250/391]	Time	0.040 (0.171)	Data	0.000 (0.136)	Loss	2.6410e-03 (2.8833e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][260/391]	Time	0.357 (0.171)	Data	0.329 (0.136)	Loss	2.4809e-03 (2.8688e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][270/391]	Time	0.041 (0.171)	Data	0.000 (0.137)	Loss	2.7665e-03 (2.8699e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][280/391]	Time	0.433 (0.171)	Data	0.404 (0.137)	Loss	2.5914e-03 (2.8610e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][290/391]	Time	0.041 (0.171)	Data	0.000 (0.136)	Loss	2.4351e-03 (2.8591e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][300/391]	Time	0.365 (0.171)	Data	0.337 (0.137)	Loss	2.6273e-03 (2.8493e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][310/391]	Time	0.041 (0.172)	Data	0.000 (0.137)	Loss	4.8451e-03 (2.8486e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][320/391]	Time	0.548 (0.172)	Data	0.520 (0.138)	Loss	2.6195e-03 (2.8586e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][330/391]	Time	0.042 (0.172)	Data	0.000 (0.138)	Loss	2.5894e-03 (2.8629e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][340/391]	Time	0.390 (0.172)	Data	0.362 (0.138)	Loss	2.2831e-03 (2.8722e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][350/391]	Time	0.041 (0.173)	Data	0.000 (0.138)	Loss	2.4506e-03 (2.8790e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][360/391]	Time	0.332 (0.173)	Data	0.304 (0.139)	Loss	2.5716e-03 (2.8803e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][370/391]	Time	0.040 (0.173)	Data	0.000 (0.138)	Loss	5.2257e-03 (2.8855e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][380/391]	Time	0.377 (0.173)	Data	0.343 (0.139)	Loss	2.4514e-03 (2.8864e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Epoch: [89][390/391]	Time	0.037 (0.173)	Data	0.000 (0.139)	Loss	2.8977e-03 (2.8924e-03)	Acc@1	100.00 (99.98)	Acc@5	100.00 (100.00)
Test: [0/40]	Time	0.404 (0.404)	Loss	2.7293e+00 (2.7293e+00)	Acc@1	41.02 (41.02)	Acc@5	71.48 (71.48)		
Test: [10/40]	Time	0.015 (0.074)	Loss	3.2318e+00 (3.2042e+00)	Acc@1	37.11 (36.61)	Acc@5	65.23 (64.13)		
Test: [20/40]	Time	0.128 (0.063)	Loss	3.5614e+00 (3.4057e+00)	Acc@1	32.81 (34.56)	Acc@5	55.86 (60.53)		
Test: [30/40]	Time	0.015 (0.057)	Loss	3.7106e+00 (3.4797e+00)	Acc@1	29.30 (34.12)	Acc@5	52.73 (58.97)		
* Acc@1 35.090 Acc@5 59.920										
Training Finished										

二、复现 Word-level Language Model 并讨论

(1) 复现训练和文本生成的过程。要求使用 Transformer 模型。提供实验截图

epoch	1		200/ 2983 batches		lr 5.00		ms/batch 13.03		loss 7.58		ppl 1956.83
epoch	1		400/ 2983 batches		lr 5.00		ms/batch 11.91		loss 6.79		ppl 891.15
epoch	1		600/ 2983 batches		lr 5.00		ms/batch 11.90		loss 6.49		ppl 660.43
epoch	1		800/ 2983 batches		lr 5.00		ms/batch 11.94		loss 6.35		ppl 574.80
epoch	1		1000/ 2983 batches		lr 5.00		ms/batch 12.57		loss 6.25		ppl 519.56
epoch	1		1200/ 2983 batches		lr 5.00		ms/batch 11.96		loss 6.22		ppl 504.81
epoch	1		1400/ 2983 batches		lr 5.00		ms/batch 12.00		loss 6.14		ppl 465.38
epoch	1		1600/ 2983 batches		lr 5.00		ms/batch 12.11		loss 6.15		ppl 470.76
epoch	1		1800/ 2983 batches		lr 5.00		ms/batch 12.05		loss 6.03		ppl 415.32
epoch	1		2000/ 2983 batches		lr 5.00		ms/batch 12.01		loss 6.02		ppl 412.79
epoch	1		2200/ 2983 batches		lr 5.00		ms/batch 11.99		loss 5.92		ppl 373.97
epoch	1		2400/ 2983 batches		lr 5.00		ms/batch 12.00		loss 5.94		ppl 378.05
epoch	1		2600/ 2983 batches		lr 5.00		ms/batch 12.00		loss 5.93		ppl 375.87
epoch	1		2800/ 2983 batches		lr 5.00		ms/batch 11.99		loss 5.84		ppl 343.72

end of epoch	1		time: 37.67s		valid loss 5.74		valid ppl 312.56
--------------	---	--	--------------	--	-----------------	--	------------------

epoch	2		200/ 2983 batches		lr 5.00		ms/batch 12.43		loss 5.80		ppl 329.07
epoch	2		400/ 2983 batches		lr 5.00		ms/batch 12.11		loss 5.76		ppl 318.44
epoch	2		600/ 2983 batches		lr 5.00		ms/batch 12.00		loss 5.62		ppl 275.87
epoch	2		800/ 2983 batches		lr 5.00		ms/batch 12.03		loss 5.62		ppl 277.15
epoch	2		1000/ 2983 batches		lr 5.00		ms/batch 12.16		loss 5.60		ppl 271.13
epoch	2		1200/ 2983 batches		lr 5.00		ms/batch 12.07		loss 5.61		ppl 274.13
epoch	2		1400/ 2983 batches		lr 5.00		ms/batch 11.96		loss 5.62		ppl 274.79
epoch	2		1600/ 2983 batches		lr 5.00		ms/batch 11.97		loss 5.66		ppl 286.68
epoch	2		1800/ 2983 batches		lr 5.00		ms/batch 11.97		loss 5.54		ppl 255.65
epoch	2		2000/ 2983 batches		lr 5.00		ms/batch 11.97		loss 5.58		ppl 264.80
epoch	2		2200/ 2983 batches		lr 5.00		ms/batch 11.98		loss 5.48		ppl 239.90
epoch	2		2400/ 2983 batches		lr 5.00		ms/batch 11.99		loss 5.52		ppl 249.24
epoch	2		2600/ 2983 batches		lr 5.00		ms/batch 12.00		loss 5.53		ppl 251.04
epoch	2		2800/ 2983 batches		lr 5.00		ms/batch 12.00		loss 5.45		ppl 233.63

end of epoch	2		time: 37.57s		valid loss 5.54		valid ppl 255.67
--------------	---	--	--------------	--	-----------------	--	------------------

	end of epoch		4		time: 38.32s			valid loss		5.41		valid ppl		222.56				

	epoch	5		200/ 2983 batches			lr	5.00		ms/batch		12.00		loss	5.08		ppl	161.18
	epoch	5		400/ 2983 batches			lr	5.00		ms/batch		11.98		loss	5.09		ppl	163.10
	epoch	5		600/ 2983 batches			lr	5.00		ms/batch		12.04		loss	4.92		ppl	136.78
	epoch	5		800/ 2983 batches			lr	5.00		ms/batch		12.06		loss	4.97		ppl	143.71
	epoch	5		1000/ 2983 batches			lr	5.00		ms/batch		12.10		loss	4.97		ppl	143.40
	epoch	5		1200/ 2983 batches			lr	5.00		ms/batch		12.16		loss	4.99		ppl	146.84
	epoch	5		1400/ 2983 batches			lr	5.00		ms/batch		12.11		loss	5.03		ppl	153.10
	epoch	5		1600/ 2983 batches			lr	5.00		ms/batch		12.09		loss	5.08		ppl	160.80
	epoch	5		1800/ 2983 batches			lr	5.00		ms/batch		12.08		loss	4.98		ppl	145.04
	epoch	5		2000/ 2983 batches			lr	5.00		ms/batch		12.07		loss	5.02		ppl	151.37
	epoch	5		2200/ 2983 batches			lr	5.00		ms/batch		12.10		loss	4.90		ppl	134.87
	epoch	5		2400/ 2983 batches			lr	5.00		ms/batch		12.06		loss	4.96		ppl	142.50
	epoch	5		2600/ 2983 batches			lr	5.00		ms/batch		12.05		loss	4.98		ppl	145.54
	epoch	5		2800/ 2983 batches			lr	5.00		ms/batch		12.07		loss	4.92		ppl	136.58

	end of epoch		5		time: 37.58s			valid loss		5.38		valid ppl		216.53				

	epoch	6		200/ 2983 batches			lr	5.00		ms/batch		12.42		loss	4.95		ppl	141.23
	epoch	6		400/ 2983 batches			lr	5.00		ms/batch		12.02		loss	4.97		ppl	143.79
	epoch	6		600/ 2983 batches			lr	5.00		ms/batch		12.11		loss	4.79		ppl	120.69
	epoch	6		800/ 2983 batches			lr	5.00		ms/batch		12.09		loss	4.85		ppl	127.36
	epoch	6		1000/ 2983 batches			lr	5.00		ms/batch		12.02		loss	4.85		ppl	127.34
	epoch	6		1200/ 2983 batches			lr	5.00		ms/batch		12.04		loss	4.87		ppl	129.95
	epoch	6		1400/ 2983 batches			lr	5.00		ms/batch		12.07		loss	4.92		ppl	136.44
	epoch	6		1600/ 2983 batches			lr	5.00		ms/batch		12.08		loss	4.96		ppl	143.11
	epoch	6		1800/ 2983 batches			lr	5.00		ms/batch		12.06		loss	4.86		ppl	129.65
	epoch	6		2000/ 2983 batches			lr	5.00		ms/batch		12.05		loss	4.90		ppl	134.91
	epoch	6		2200/ 2983 batches			lr	5.00		ms/batch		12.10		loss	4.79		ppl	120.13
	epoch	6		2400/ 2983 batches			lr	5.00		ms/batch		12.03		loss	4.84		ppl	126.71
	epoch	6		2600/ 2983 batches			lr	5.00		ms/batch		12.09		loss	4.87		ppl	130.33
	epoch	6		2800/ 2983 batches			lr	5.00		ms/batch		12.06		loss	4.81		ppl	122.12

	end of epoch		6		time: 37.64s			valid loss		5.37		valid ppl		214.04				

=====																		
	End of training			test loss		5.28		test ppl		195.58								
=====																		

(2) Transformer 和 CNN 在捕捉上下文依赖上有什么差异？

神经网络结构：Transformer 是基于自注意力机制的序列到序列模型，其中包含 Encoder 和 Decoder 两部分，每个部分包含多个层，相邻层之间使用残差连接和 Layer Normalization 进行连接，利用自注意力机制来获取输入序列的全局信息。而 CNN 则是一种卷积神经网络，其主要特点是使用卷积操作来对输入进行特征提取，对于输入数据的每个不同区域使用不同的权重进行加权汇聚，从而捕捉空间上的局部相关性。

应用场景：Transformer 主要用于序列到序列模型训练任务，例如机器翻译、语音识别等。而 CNN 主要用于计算机视觉任务，如图像分类、物体检测等任务。

模型复杂度：Transformer 由于需要考虑全局信息，因此模型参数较多。CNN 则只需要考虑空间相邻区域的信息，因此模型参数较少。

捕获依赖方式：Transformer 通过自注意力机制来学习序列中不同位置之间的依赖。对于每个位置来说，通过乘以一个权重矩阵并加和得到新的特征向量表示。这种方式可以捕获不同位置之间的长距离依赖。而 CNN 主要是在局部区域内进行卷积操作，对于不同的局部区域使用不同的卷积核进行加权汇聚，从而捕捉局部相关性。